

**Jin Gu**

Department of Automation, Tsinghua University

Email: [jgu@tsinghua.edu.cn](mailto:jgu@tsinghua.edu.cn)

Phone: (010) 62794294-866

# Chapter 10 Parameter Learning with Complete Data

2021 Fall

Jin Gu (古槿)

# Move to *Learning*

- Representation

$$P \Leftrightarrow \{P, G\}$$

- Parents  $\rightarrow$  child structures & cliques
- Gaussian models & exponential families

- Inference

- Particle-based inference
- Inference as optimization

$$P(Y \mid E = e, \theta)$$

- Learning

$$\max_{\theta} P(\mathbf{x}[1], \mathbf{x}[2], \dots \mid \theta)$$

$$P(\theta \mid \mathbf{x}[1], \mathbf{x}[2], \dots)$$

# Three Lectures

- **Parameter** learning with complete data
  - Learning with **hidden** models or missing values
  - **Structure** learning
  - Learning with deep models & general graphs
- 
- ***Quiz II, 2021/12/20, Inference & Learning***
  - OOH: Friday 8:30-9:30
  - Anytime with TAs (make appointment by email)

# Outlines

- Learning Basics
- **Maximum Likelihood** Parameter Estimation
  - Basics for Maximum Likelihood Estimation
  - Maximum Likelihood Estimation in BNs
- **Bayesian** Parameter Estimation
  - Basics for Bayesian Estimation
  - Priors for Bayesian Estimation
  - Bayesian Estimation in BNs
- **MAP** Parameter Estimation

# Outlines

- Parameter Learning in Markov Networks
  - The likelihood of Markov networks
  - The convexity of partition function
    - Recall: stochastic gradient descent
      - Example: CRFs learning revisited
- *Approximate methods*
  - *Pseudo likelihood (cannot deal with hidden values)*
  - *Contrastive divergence*

## **Chapter 10**    Parameter Learning with Complete Data

### **Textbook1**

**Chapter 17.1-17.2** Maximum Likelihood Estimation in Bayesian Networks

**Chapter 17.3-17.4** Bayesian Learning in Bayesian Networks

**Chapter 20.1-20.3** Maximum Likelihood Estimation in Markov Networks

### **Textbook2**

**Chapter 10.4** Learning in Directed Graphical Models

**Chapter 19.5** Learning in Undirected Graphical Models

# What is Learning?

- Draw *probabilistic models* from *observed data*

- Learning as *optimization*  $\max_{\theta} P(\mathbf{x}[1], \mathbf{x}[2], \dots \mid \theta)$

- Find *an optimal probabilistic model* which can maximize a given objective function

- Learning as *probabilistic* inference

- Treat the parameters as variables and *infer their posteriors* based on observed data

$$P(\theta \mid \mathbf{x}[1], \mathbf{x}[2], \dots)$$

# Learning as Optimization

- Maximum likelihood  $\theta^* = \max_{\theta} P(\mathbf{x}[1], \mathbf{x}[2], \dots, \mathbf{x}[M] | \theta)$
- Maximum a posterior  $\theta^* = \max_{\theta} P(\theta | \mathbf{x}[1], \mathbf{x}[2], \dots, \mathbf{x}[M])$
- *Not as a probabilistic problem*
  - *As minimum error*
  - *As maximum margin*
  - *As compressive sensing*
  - ....

Define proper *loss* or *likelihood* functions *based on different principles*



# Learning as Probabilistic Inference

- Learning in a **probabilistic** framework
  - Parameter learning:  $\{x[m]\}_{m=1 \sim M} |_{\mathcal{G}} \rightarrow P(\theta | \mathcal{D})$
  - Structure learning:  $\{x[m]\}_{m=1 \sim M} \rightarrow P(\mathcal{G}, \theta | \mathcal{D})$
- Generative models
  - Learn the **joint** representation  $\tilde{P}(Y, X | \theta)$
- Discriminative models
  - Learn the **conditional** probability  $\tilde{P}(Y | X = x, \theta)$

# Learning Basics: IID Samples

- For most learning tasks, the variables are required to be “independent and identical distributed” (IID or *i.i.d.*). The observed samples are generated from those variables are called IID samples.
- IID means that the generation of a specific sample **is unrelated** to all the other samples
- *How to collect samples from Gibbs sampling?*

# Learning Basics: Avoid Overfitting

- Overfitting
  - If the *model complexity* (or the freedom of model) is much larger than the training data, we can always get “good” learning with zero *empirical risk*
  - But when testing on new data, the learned model does bad predictions
- Generalization
  - To avoid overfitting or increase the generalization of learning methods, we need to *penalize the model complexity* and *separate training/testing*

# Avoid Overfitting (Theoretical)

- The regularization term in loss function
  - Bayesian learning
  - Statistical learning theory: structural risk
  - Sparsity
  - Low-rank
  - .....

**Model complexity** should fit  
the complexity of data!

# Avoid Overfitting (Empirical)

- Cross-validation
  - LOOCV (leave one out cross-validation)
  - N-fold cross-validation
- 0.632 bootstrapping ( $0.632 = 1 - e^{-1}$ )

Training performance should be similar  
with testing performance!

# 0.632 Bootstrapping

- For  $M$  samples,  $x[1], \dots, x[M]$
- Randomly generated  $M$  new samples by with-replacement sampling: each time randomly choose a sample  $x[k]$  and then put it back to the sample pool
- On average,  $0.632 * M$  samples will be selected into the training dataset
- The left  $0.368 * M$  samples will be used as the testing dataset

## 0.632 Bootstrapping

- Then, learn your model on the training dataset and test its performance on the testing data
- Repeat above steps many times
- Check whether the performances are significant lower in testing than training (if yes, your learning is over-fitting)
- If similar, you get the estimated performance  $0.368 * Perform_{Train} + 0.632 * Perform_{Test}$

# Parameter Estimation



- For a set of samples, they are correlated if the distribution parameters are not given or observed

$$P(x[1], x[2], \dots, x[M]) \neq P(x[1]) \cdots P(x[M])$$

- If we assume the distribution parameters are given, all the samples are *i.i.d.*

$$P(x[1], \dots, x[M]|\theta) = P(x[1]|\theta) \cdots P(x[M]|\theta)$$

*i.i.d* assumption!





# Maximum Likelihood Parameter Estimation

- What is *likelihood*?
  - The probability or confidence for parameter assignment, given a number of data
  - Log likelihood is commonly used for better calculation
    - $l(\theta; \mathcal{D}) \propto \log \tilde{P}(\mathcal{D}|\theta) = \sum_i \log \phi_i(x[i]; \theta)$
- *MLE*: the parameter estimation is to find the optimal parameter assignment  $\theta^*$  which can maximize the likelihood (optimization)

# MLE: Maximum Likelihood Estimation

- We throw a five cent coin  $M$  times with  $M[T]$  text-face up and  $M[L]$  as lotus-face up. Please find  $\theta^*$  of the coin ( $P(T) = \theta$ ) which is most likely to generate that result.

- The log-likelihood

$$l(\theta: X) = \log(\theta^{M[T]}(1 - \theta)^{M[L]}).$$

- Set the derivative of the log-likelihood as 0

$$\frac{M[T]}{\theta^*} - \frac{M[L]}{1 - \theta^*} = 0 \rightarrow \theta^* = \frac{M[T]}{M[T] + M[L]}.$$

# Maximum Likelihood Estimation

- In many situations, optimization algorithms should be used to maximize the likelihood function (commonly in log-linear format)
  - For complex likelihood function, it is hard to get the close form of its derivative
  - For incremental or updating learning, we want to improve the likelihood sample by sample
  - Regularization terms can be added to control the model complexity (such as LASSO)

# Gradient-Based Methods

- Gradient descent/ascent is a common strategy to find the optimal parameter setting. Simply update the model along the gradient:

$$\boldsymbol{\theta}^{(k+1)} = \boldsymbol{\theta}^{(k)} + \lambda^{(k)} \nabla l(\boldsymbol{\theta}; \mathbf{x})$$

- If the likelihood function  $l(\boldsymbol{\theta}; \mathbf{x})$  is **convex**, the gradient ascent will **converge to the optimal**
- In most cases  $\theta_i$  is independent with each other, you can update them one by one using partial derivative on  $\theta_i$

# Stochastic Gradient Ascent

- If the observed samples are *i.i.d.*, the likelihood function are the sum across each sample

$$l(\boldsymbol{\theta}; \boldsymbol{x}) = \sum_{m=1}^M l_m(\boldsymbol{\theta}; \boldsymbol{x}[m]).$$

- In many situations, the computational cost is huge for the sum
- Stochastic gradient ascent **updates parameters based on one sample each time**


$$\boldsymbol{\theta}^{(k+1)} = \boldsymbol{\theta}^{(k)} + \lambda^{(k)} \nabla l_{m_k}(\boldsymbol{\theta}; \boldsymbol{x}[m_k]).$$

$$\boldsymbol{\theta}^{(k+1)} = \boldsymbol{\theta}^{(k)} + \lambda^{(k)} \nabla l_{m_k}(\boldsymbol{\theta}; \mathbf{x}[m_k])$$

- The order of the input samples are randomly given in each run
- The input samples can be used repeatedly (each time with a different random order) until the optimization process converges

Comments: to avoid *local minimal*, heuristic methods should be used, such as simulated annealing and genetic algorithm

# Logistic Regression Gradient



- Logistic regression

$$p(y=1 | \bar{x}) = \frac{e^z}{1+e^z}, \quad z = -\sum_j \beta_j x_j$$

- If  $y = 0$  the partial derivative of  $\log(1 - p)$  is

$$\frac{\partial}{\partial \beta_j} \log(1 - p) = \frac{\partial}{\partial \beta_j} \left( \log \left( \frac{1}{1+e^z} \right) \right) = p x_j$$

- If  $y = 1$  the partial derivation of  $\log(p)$  is

$$\frac{\partial}{\partial \beta_j} \log(p) = \frac{\partial}{\partial \beta_j} \left( \log \left( \frac{e^z}{1+e^z} \right) \right) = \frac{\partial}{\partial \beta_j} \left( \log(e^z) + \log \left( \frac{1}{1+e^z} \right) \right) = (p - 1) x_j$$

# Logistic Regression Gradient

- So the partial derivative can be written as

$$\sum_{i:y=0} p_i x_{ij} + \sum_{i:y=1} (p_i - 1) x_{ij} = \sum_i (p_i - y_i) x_{ij}$$

- Set the partial derivation as 0, so we get

$$\sum_i y_i x_{ij} = \sum_i p_i x_{ij}$$

- This equation can be used to check the correctness of a trained model



# Logistic Regression Gradient

For **any single training sample**, the partial derivative of the log likelihood is

$$\frac{\partial}{\partial \beta_j} \log L(x, y; \beta) = (y - p) x_j$$

- Learning algorithm:
  - The parameters are given initial values  $\beta^{(0)}$
  - Calculate  $p^{(0)}$  according to the regression
  - Then, update the parameters based on the partial derivative

$$\beta_j^{(k+1)} := \beta_j^{(k)} + \lambda^{(k)} (y - p^{(k)}) x_j$$

Learning rate  $\lambda$  should be decreased as the iterative steps increase

# MLE in Bayesian Networks

- Global likelihood decomposition


- Recall:  $P(X) = \prod_i P(X_i | Pa_{X_i})$

$$L(\theta: X) = \prod_m P(x[m]; \theta) = \prod_i \prod_m P(x_i[m] | Pa_{x_i[m]}; \theta_{x_i | Pa_{x_i}})$$

- We get the **local likelihood** function

$$L_i(\theta: X) = \prod_m P(x_i[m] | Pa_{x_i[m]}; \theta_{x_i | Pa_{x_i}})$$

- According to the decomposition, we can **maximize each local likelihood independently** and combine them as an MLE solution



# MLE in Bayesian Networks

- For a local likelihood, only take the data scopes of the variables  $X_i$  and its parents  $Pa_{X_i}$
- Learning as table CPDs
  - Count the frequency as the parameter
$$\theta_{X_i=x_i|Pa_{X_i}=Pa_{x_i}} = M[x_i, Pa_{x_i}] / M[Pa_{x_i}]$$
- Learning as generalized linear models
  - The same approach as logistic regression above



*The learning process is **totally independent!***

If the network structure is sparse, the learning algorithm can work very efficiently

# Bayesian Parameter Estimation

- MLE is for *point estimation*, but not for *prediction* or *distribution estimation*
- MLE is maximizing the joint probability of the observed data for finding the optimal parameter, but it *cannot use the prior knowledge or constraint* of the parameter
- We can do estimation using Bayes' rule

$$P(\theta|X) = \frac{P(X|\theta)P(\theta)}{\int P(X|\theta)P(\theta)d\theta}$$



- We get the posterior distribution of the parameter rather than a point estimation
- In most cases, the Bayesian model is used to do prediction (on new data)

$$P(X^{n+1}|X^1, \dots, X^n) = \frac{1}{P(X^1, \dots, X^n)} \int P(X^{n+1}|\theta)P(X^1, \dots, X^n|\theta)P(\theta)d\theta$$

- Recall the first homework: we assume  $\theta$  is set as 0.5, and we observed 63 “upward” out of 100 random throws. Please show the probability that we get “upward” for the 101-th throw.

# Priors for Parameters

- The priors of parameters

$$\frac{1}{P(X^1, \dots, X^n)} \int P(X^{n+1} | \theta) P(X^1, \dots, X^n | \theta) \mathbf{P}(\boldsymbol{\theta}) d\theta$$


- Use *uniform* distribution if no prior
- The *Beta* distribution is more commonly used

$$\theta \sim \text{Beta}(\alpha_1, \alpha_0) \text{ if } p(\theta) = \gamma \theta^{\alpha_1-1} (1 - \theta)^{\alpha_0-1}$$


- $\gamma = \frac{\Gamma(\alpha_1 + \alpha_0)}{\Gamma(\alpha_1)\Gamma(\alpha_0)}$  is a normalizing factor
- $\Gamma(x) = \int_0^\infty t^{x-1} e^{-t} dt$ ,  $\Gamma(n) = (n-1)!$  for integer

$$p(\theta) = \gamma \theta^{\alpha_1 - 1} (1 - \theta)^{\alpha_0 - 1}$$


# Priors: Beta Distribution




*Beta(1,1)*




*Beta(2,2)*




*Beta(10,10)*



*Beta(3,2)*



*Beta(15,10)*



*Beta(0.5,0.5)*

# Priors: Beta Distribution

- For parameter distribution estimation

$$P(\theta | X) = \frac{P(X | \theta) P(\theta)}{\int P(X | \theta) P(\theta) d\theta}$$

All calculations are **in the same family**. We say that **Beta distribution is conjugate**.

$$= \frac{\theta^{M[1]} (1-\theta)^{M[0]} \gamma \theta^{\alpha_1-1} (1-\theta)^{\alpha_0-1}}{\int \theta^{M[1]} (1-\theta)^{M[0]} \gamma \theta^{\alpha_1-1} (1-\theta)^{\alpha_0-1} d\theta} = \text{Beta}(M[1] + \alpha_1, M[0] + \alpha_0)$$



# Priors: Beta Distribution

$$(uv)' = u'v + uv' \Rightarrow \int_a^b uv' = uv \Big|_a^b - \int_a^b u'v$$

- For prediction

$$\begin{aligned} P(x[m+1]=1 | X) &= \int P(x[m+1] | \theta, X) P(\theta | X) d\theta \\ &= \int P(x[m+1] | \theta) P(\theta | X) d\theta \end{aligned}$$

$$P(\theta | X) \propto P(X | \theta) P(\theta) \propto \theta^{M[1]} (1-\theta)^{M[0]} \theta^{\alpha_1-1} (1-\theta)^{\alpha_0-1}$$

$$\therefore P(x[m+1]=1 | X) = \int \theta P(\theta | X) d\theta = \frac{M[1] + \alpha_1}{M + \alpha}$$

- Special case: uniform prior

$$P(x[m+1]=1 | X) = \frac{M[1] + 1}{M + 2}$$

– We have *Laplace's correction* for the prediction

*How about doing prediction based on the MLE parameter?*

# Priors: *Dirichlet* Distribution

- *Beta distribution* can only be used for binomial distribution (one independent parameter) with constraint  $\theta_1 + \theta_2 = 1$
- *Dirichlet distribution* is an extension for multinomial distribution  $\sum_k \theta_k = 1$
- The likelihood function for multinomial

$$L(\theta; \mathcal{D}) = \prod_k \theta_k^{M[k]}$$

- The *Dirichlet* distribution

$$\theta \sim \text{Dirichlet}(\alpha_k) \text{ if } P(\theta) \propto \prod_k \theta_k^{\alpha_k - 1}$$

# Priors: *Dirichlet* Distribution

- *Dirichlet* distribution has good properties
  - It is conjugate
  - Its posterior is simple  $\alpha_k^* = M[k] + \alpha_k$
- So, we can quickly get the prediction

$$P(x[M + 1] = k) = \frac{M[k] + \alpha_k}{M + \alpha}$$


- Let  $\theta'_k$  is the prior expectation for  $x[M + 1] = k$ . Surely,  $\theta'_k = \alpha_k / \alpha$

$$\frac{M[k] + \alpha_k}{M + \alpha} = \frac{\alpha}{M + \alpha} \theta'_k + \frac{M}{M + \alpha} \frac{M[k]}{M}$$

We can treat the posterior as the **weighted average** between **prior** and **MLE estimation**

# Bayesian Estimation in BNs

- Meta-network for IID samples:



- So we can do global decomposition, if  $P(\theta)$  satisfies the global independence:

$$P(\theta|\mathcal{D}) = \prod_i P\left(\theta_{X_i|Pa_{X_i}}|\mathcal{D}\right)$$

# Bayesian Estimation in BNs

- Then we can further do local decomposition

$$\prod_i P\left(\theta_{X_i|Pa_{X_i}}|\mathcal{D}\right) = \prod_i \prod_{pa_{X_i}} P\left(\theta_{X_i|pa_{X_i}}|\mathcal{D}\right)$$

- If we use the *Dirichlet* distribution as prior, the Bayesian prediction should be

$$P\left(X_i[M+1] = x_{ij} | Pa_{X_i}[M+1] = u, \mathcal{D}\right) = \frac{\alpha_{x_{ij}|u+M}[x_{ij}, u]}{\sum_j \alpha_{x_{ij}|u+M}[u]}$$

# Bayesian Estimation in BNs

How to choose the parameters in prior?

- *BDe prior*: use another distribution  $P'(X)$  to get the priors
  - Usually the same structure as  $\mathcal{G}$ , but not required
  - First set the total confidence  $\alpha$  of the prior, which represent the *pseudo* counts in predictions
    - Set  $\alpha_{x_{ij}|pa_{X_i}} = \alpha \cdot P'(X_i = x_{ij}, Pa_i = pa_{X_i})$
    - Standard inferences are needed to calculate the marginal distribution in  $P'$

# MAP Parameter Estimation

- If we can get a full solution of the Bayesian estimation, MAP estimation is defined as

$$\tilde{\theta} = \arg \max_{\theta} \log P(\theta|\mathcal{D})$$

- For example, for binomial distribution with Beta prior  $\alpha_1 = \alpha_0 = 1$ :  $P(\theta|\mathcal{D}) \propto \theta^{M[1]}(1 - \theta)^{M[0]}$ . The MAP estimation is the same as the MLE.
- *Optimization methods* is usually used for MAP estimation if the full solution is hard to calculate

*Maximum a posterior* (极大后验)

# MAP Parameter Estimation

- When we have a large amount of data, the MAP estimation is mainly affected by training data

$$\tilde{\theta} = \arg \max_{\theta} (\log P(\mathcal{D}|\theta) + \log P(\theta))$$

- The likelihood function  $\log P(\mathcal{D}|\theta)$  will dominate
  - $\log P(\theta)$  can be regarded as *regularization*
- The posterior for prediction is often sharply peaked around its MAP estimation

$$P(X|\mathcal{D}) = \int P(X|\theta)P(\theta|\mathcal{D})d\theta \approx P(X|\tilde{\theta})$$

- The **shape** of Bayesian distribution of  $\tilde{\theta}$  is **sharp**



# Comments

- Bayesian estimation can **control the generalization ability** when **the training data are not enough** relative to the model complexity (many assignments will get zero probability for MLE)
- Using Bayesian estimation to do predictions can be regarded as a **model averaging** over all possible parameter settings
- When the training data are increasing, the weight of the prior is reducing. The resulting estimation will be mainly determined by the data
- When the number of data is large, the shape of the posterior will be sharply peaked around MAP

# The Likelihood of Markov Networks

- For discretized random variables

$$P(X) = \frac{1}{Z} \exp(\sum_i \theta_i f_i(C_i))$$

$$l(\theta: D) = \sum_i \theta_i \sum_m f_i(x_i[m]) - M \ln Z$$

$$\frac{1}{M} l(\theta: D) = \sum_i \theta_i E_D(f_i(x_i)) - \ln Z$$

- **Bad: the second term couples all factors!**

$$Z = \sum_{c_i} \exp(\sum_i \theta_i f_i(c_i))$$

# The Convexity of Partition Function

- Convexity

$$f(\alpha x + (1 - \alpha)y) \leq \alpha f(x) + (1 - \alpha)f(y)$$

- Hessian matrix is semi-positive  $\Rightarrow$  convexity
- Hessian matrix is the matrix of the secondary partial derivatives

$$H(f) = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \cdots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}.$$

# The First Derivative of $Z$

- The first derivative of  $Z$  is the expectation of each feature function

$$\begin{aligned}\frac{\partial}{\partial \theta_i} \ln Z(\boldsymbol{\theta}) &= \frac{1}{Z(\boldsymbol{\theta})} \sum_{\xi} \frac{\partial}{\partial \theta_i} \exp \left\{ \sum_j \theta_j f_j(\xi) \right\} \\ &= \frac{1}{Z(\boldsymbol{\theta})} \sum_{\xi} f_i(\xi) \exp \left\{ \sum_j \theta_j f_j(\xi) \right\} \\ &= \mathbf{E}_{\boldsymbol{\theta}}[f_i].\end{aligned}$$

# The Secondary Derivative of Z

$$\begin{aligned}
 \frac{\partial^2}{\partial \theta_j \partial \theta_i} \ln Z(\theta) &= \frac{\partial}{\partial \theta_j} \left[ \frac{1}{Z(\theta)} \sum_{\xi} f_i(\xi) \exp \left\{ \sum_k \theta_k f_k(\xi) \right\} \right] \\
 &= -\frac{1}{Z(\theta)^2} \left( \frac{\partial}{\partial \theta_j} Z(\theta) \right) \sum_{\xi} f_i(\xi) \exp \left\{ \sum_k \theta_k f_k(\xi) \right\} \\
 &\quad + \frac{1}{Z(\theta)} \sum_{\xi} f_i(\xi) f_j(\xi) \exp \left\{ \sum_k \theta_k f_k(\xi) \right\} \\
 &= -\frac{1}{Z(\theta)^2} Z(\theta) E_{\theta}[f_j] \sum_{\xi} f_i(\xi) \tilde{P}(\xi : \theta) \\
 &\quad + \frac{1}{Z(\theta)} \sum_{\xi} f_i(\xi) f_j(\xi) \tilde{P}(\xi : \theta) \\
 &= -E_{\theta}[f_j] \sum_{\xi} f_i(\xi) P(\xi : \theta) \\
 &\quad + \sum_{\xi} f_i(\xi) f_j(\xi) P(\xi : \theta) \\
 &= E_{\theta}[f_i f_j] - E_{\theta}[f_i] E_{\theta}[f_j] \\
 &= \text{Cov}_{\theta}[f_i; f_j].
 \end{aligned}$$

The secondary derivative is the *covariance matrix* between different factors

# The Convexity of Partition Function

- Hessian matrix of the partition function is the covariance matrix of the feature functions
- As we known, *covariance matrix* is always positive semi-definite
- The partition function is **convex**. So, we can implement efficient *gradient ascent algorithm* to find the global optimal solution for the likelihood function!

# Recall: Gradient Ascent

- Gradient descent is a common strategy to find the optimal parameter setting. Simply update the model along the gradient:

$$\boldsymbol{\theta}^{(k+1)} = \boldsymbol{\theta}^{(k)} + \lambda^{(k)} \nabla l(\boldsymbol{\theta}; \mathbf{x})$$

- If the likelihood function  $l(\boldsymbol{\theta}; \mathbf{x})$  is convex, the gradient ascent will converge to the optimal
- In most cases  $\theta_i$  is independent with each other, you can update them one by one using partial derivative on  $\theta_i$

# Stochastic Gradient Ascent

- If the observed samples are *i.i.d.*, the likelihood function are the sum across each sample


$$l(\boldsymbol{\theta}; \mathbf{x}) = \sum_{m=1}^M l_m(\boldsymbol{\theta}; \mathbf{x}[m])$$

- In many situations, the computational cost is huge for the sum
- Stochastic gradient ascent updates parameters based on one sample each time

$$\boldsymbol{\theta}^{(k+1)} = \boldsymbol{\theta}^{(k)} + \lambda^{(k)} \nabla l_{m_k}(\boldsymbol{\theta}; \mathbf{x}[m_k])$$



- The order of the input samples are randomly given in each run
- The input samples can be used repeatedly (each time with a different random order) until the optimization process converges
- Comments: to avoid local minimal, heuristic methods should be used, such as simulated annealing and genetic algorithm




$$x: \alpha_1 = \alpha_2 = \alpha_3 = \alpha$$

$$y: \beta$$

$$xx: w_{12} = w_{13} = w_{23} = w$$

$$xy: h_{1y} = h_{2y} = h_{3y} = h$$

# Linear-Chain CRFs (Revisited)



$$P(Y | X; \theta) \propto \exp \left( \sum_{i=1}^n \sum_m w_m f_m(Y_{i-1}, Y_i, X, i) \right)$$

$$P(\bar{y} | \bar{x}; \theta) = \frac{1}{Z(w, \bar{x})} \exp \left( \sum_{i=1}^n \sum_m w_m f_m(y_{i-1}, y_i, \bar{x}, i) \right)$$

$$Z(w, \bar{x}) = \sum_{\bar{y}} \exp \left( \sum_{i=1}^n \sum_m w_m f_m(y_{i-1}, y_i, \bar{x}, i) \right)$$

# Stochastic Gradient Ascent for CRFs

- Calculate the partial derivative of each  $w_m$
- Let  $F_m(\mathbf{x}, \mathbf{y}) = \sum_i f_m(y_{i-1}, y_i, \mathbf{x}, i)$  (sum across the sequence for the  $m$ -th feature)
- The derivative of the log likelihood is

$$\frac{\partial}{\partial w_m} \log p = F_m(\bar{x}, \bar{y}) - \sum_{\bar{s}} F_m(\bar{x}, \bar{s}) p(\bar{s} | \bar{x}; w)$$

First item is the number of that the feature is positive, **given the training label sequence**

Second item is the **expected number** of that the feature is positive **across all possible labels**

$$P(\bar{y} | \bar{x}; w) = \frac{1}{Z(w, \bar{x})} \exp\left(\sum_m w_m F_m(\bar{x}, \bar{y})\right)$$
$$Z(w, \bar{x}) = \sum_{\bar{y}} \exp\left(\sum_m w_m F_m(\bar{x}, \bar{y})\right)$$

# Stochastic Gradient Ascent for CRFs

- Update the parameters according to its partial derivative (the same as logistic regression)

$$w_m^{(k+1)} := w_m^{(k)} + \lambda^{(k)} \left( F_m(\bar{x}, \bar{y}) - \sum_{\bar{s}} F_m(\bar{x}, \bar{s}) p(\bar{s} | \bar{x}; w^{(k)}) \right)$$

- The **computational cost is huge** for the second item, especially for the training sequence is very long and the scope of the possible labels is very large

# The Collins Perceptron

- Given the initial weight  $w^{(0)}$ , we can calculate the optimal sequence  $\hat{y}^{(0)}$
- Then we update  $w$  using the formula below

$$w_m^{(k+1)} := w_m^{(k)} + \lambda^{(k)} \left( F_m(\bar{x}, \bar{y}) - F_m(\bar{x}, \hat{y}^*) \right)$$

- Collins. **Discriminative training methods for hidden Markov models: Theory and experiments with the perceptron algorithm.** *EMNLP 2002.*

# The End of Chapter 10