

Jin Gu

Department of Automation, Tsinghua University

Email: jgu@tsinghua.edu.cn

Phone: (010) 62794294-866

Chapter 11 Learning with Incomplete Data

2021 Fall

Jin Gu (古槿)

第二次随堂测验和课程研讨

- 第二次随堂测验

- 第15周周一（12月20日）上课时间

- 小测验范围：推断与学习

- 重点：似然加权采样、重要性采样、吉布斯采样、平均场变分法；MLE、贝叶斯参数学习、随机梯度优化、不完全观测学习、BIC评分

- 考核方式：开卷（严禁讨论）

第二次随堂测验和课程研讨

- 第二次课程研讨
 - 第14周周二（12月14日）晚19:00-21:30
 - 后三次作业讨论
 - 课程论文讲解
 - 课程论文截止时间17周周五（1月7日）
- 请及时填写学生评教

Outlines

- Some Examples
 - Variables are *Not Detectable*
 - Hidden Markov Revisited
 - Mixture Models
 - Latent Linear Models
 - *Missing Values* and *Data Outliers*
- General Principles and Methods
 - General Principles
 - Expectation Maximization (EM)
 - MCMC Sampling

Chapter 11 Learning with Incomplete Data

Textbook1

Chapter 19.2.1 Gradient Ascent

Chapter 19.2.2.1-19.2.2.4, 19.2.2.6 Expectation Maximization (EM)

Chapter 20.1-20.3 Maximum Likelihood Estimation in Markov Networks

Textbook2

Chapter 11 Mixture models and the EM algorithm

Chapter 12 Latent Linear Models

Chapter 27* Latent variable models for discrete data


* Advanced Readings

The Reasons of Incomplete Data

- Variables are not detectable (hidden variables)
 - Variables cannot be observed
 - Variables only in concept
- Missing values and data outliers
 - The systems miss some observations
 - A few exceptional data points

Hidden Markov Models Revisited

- The state variables are not observable



$$\theta = \left\{ \begin{array}{l} T = \{t_{i,j}, \quad i, j = 1 \dots N\}, \\ E = \{e_{i,j}, \quad i = 1 \dots N, j = 1 \dots K\}, \\ \pi = \{\pi_i, \quad i = 1 \dots N\} \end{array} \right\}$$


$$X = \{x_t, \quad t = 1, \dots, T\}$$

Mixture Models

- The data are randomly drawn from two or more different distributions (or classes), but we do not know the origin of each data point
- Gaussian Mixture Models (GMMs)

$$P(X = k) = \alpha_k, \quad \sum_k \alpha_k = 1, \alpha_k > 0$$

$$p(Y|X = k) = N(\mu_k, \sigma_k^2)$$



Latent Linear Models

- Traditional linear regression (all observed)

$$Y = b + BX$$


- The observation variables X can be explained by linear combinations of a set of shared latent variables Z

$$X_i = \beta_{i,0} + \sum_{j=1}^k \beta_{i,j} z_j + \varepsilon \Leftrightarrow X = b + BZ$$

- Extended to partial least square models

$$Y = b_Y + W_Y Z_S$$

$$X = b_X + W_X Z_S + B_X Z_X$$



Learning Gaussian Mixture Models

- Learning task under the framework of MLE:

$$\arg \max_{\theta} p(\mathcal{D}|\theta) \quad \mathcal{D} = \{y[1], \dots, y[M]\}$$

- *Task Assupmtion:*

- X (label variable) is un-observed (hidden)
- The model parameters are unknown




Learning Gaussian Mixture Models

- *If X is observed*
- The learning is simple (MLE)
 - Complete data: $\mathcal{D}_C = \{(x[i], y[i])\}_{i=1 \dots M}$

$$\pi_k^* = \frac{M[x=k]}{M}$$

- For k -th mixture Gaussian component

- $\mu_k^* = \frac{1}{M[x=k]} \sum_m y[m] \mid x[m]=k$
- $\Sigma_k^* = \frac{1}{M[x=k]} \sum_m (y[m] - \mu_k^*)(y[m] - \mu_k^*)^T \mid x[m]=k$



Learning Gaussian Mixture Models

- *If the parameters are given*
- For partially observed data, we can easily infer **the posterior** of the label variable X


$$Q(x = k) = P(x = k|y, \theta) = \frac{p(y|x=k, \theta)P(x=k|\theta)}{\sum_{k=1}^K p(y|x=k, \theta)P(x=k|\theta)}$$

$$Q(x[m] = k) = \frac{\pi_k^{(t)} N_k^{(t)}(y[m])}{\sum_{k=1}^K \pi_k^{(t)} N_k^{(t)}(y[m])}$$

$$N_k(y) = \frac{1}{\sqrt{|2\pi\Sigma|}} e^{-\frac{1}{2}(y-\mu_k)^T \Sigma^{-1}(y-\mu_k)}$$

Learning Gaussian Mixture Models


- *Problem*
 - *The data are incomplete (hidden variables)*
 - *The parameters are unknown*



How to deal with this problem?

Learning Gaussian Mixture Models

- Possible solution
 - Maximize the marginal over observed variables
 - $\max_{\theta} P(D_{obs}|\theta) = \max_{\theta} \sum_X P(D_{obs}, X|\theta)$
 - Iterative algorithm
 - Set initial parameters (randomly if no prior)
 - Infer the posterior for each data point $P(x[m]|y[m], \theta)$
 - Update the parameters by a modified MLE
 - Key: use Expectation to fill the hidden variable



Expectation Maximization

Learning Gaussian Mixture Models

- *Randomly initialize the parameters*
- For partially observed data, we can easily infer **the posterior** of the label variable X

$$Q(x = k) = P(x = k|y, \theta) = \frac{p(y|x=k, \theta)P(x=k|\theta)}{\sum_{k=1}^K p(y|x=k, \theta)P(x=k|\theta)}$$

$$Q(x[m] = k) = \frac{\pi_k^{(t)} N_k^{(t)}(y[m])}{\sum_{k=1}^K \pi_k^{(t)} N_k^{(t)}(y[m])}$$

$$N_k(y) = \frac{1}{\sqrt{|2\pi\Sigma|}} e^{-\frac{1}{2}(y-\mu_k)^T \Sigma^{-1}(y-\mu_k)}$$

Learning Gaussian Mixture Models

- Parameter updating
 - The parameter π can be updated as (MLE)

$$\pi_k^{(t+1)} = \frac{1}{M} \sum_{m=1}^M Q^{(t)}(x[m] = k)$$

Compared to complete data: $\pi_k^ = \frac{1}{M} M[x = k]$*

Treat the expectation *as a weight* for the sample: a sample is divided into small pieces

Learning Gaussian Mixture Models

- Parameter updating (continued)
 - Accordingly, the parameters of $p(y|x = k)$ can also be updated based on $Q^{(t)}(x[m] = k)$

$$\mu_k^{(t+1)} = f(y, Q^{(t)}) = \frac{\sum_{m=1}^M Q^{(t)}(x[m]=k) y[m]}{\sum_{m=1}^M Q^{(t)}(x[m]=k)}$$

$$\Sigma_k^{(t+1)} = f\left(y, Q^{(t)}, \mu_k^{(t+1)}\right) = \frac{\sum_{m=1}^M Q^{(t)}(x[m]=k) (y[m] - \mu_k^{(t+1)}) (y[m] - \mu_k^{(t+1)})^T}{\sum_{m=1}^M Q^{(t)}(x[m]=k)}$$

Treat the expectation *as a weight* for the sample: a sample is divided into small pieces

Learning Gaussian Mixture Models

- Do inference after parameter updating
 - Based on the updated parameters, we can easily calculate the update posterior $Q^{(t+1)}(x[m] = k)$

$$Q(x = k) = P(x = k|y, \theta) = \frac{p(y|x=k, \theta)P(x=k|\theta)}{\sum_{k=1}^K p(y|x=k, \theta)P(x=k|\theta)}$$

$$Q(x[m] = k) = \frac{\pi_k^{(t)} N_k^{(t)}(y[m])}{\sum_{k=1}^K \pi_k^{(t)} N_k^{(t)}(y[m])}$$

Comments

$$\textit{Aim}: \arg \max_{\theta} p(\mathcal{D}|\theta)$$

- The key problem for GMM learning is that **both** *the component labels* of the samples and *the model parameters* are unknown
- We can *iteratively* update the component labels based on *posteriors* (inference) and the model parameters based on *MLE* (learning)

Q1: will this iterative process converge?

Q2: If converged, can it get the unbiased estimation?

Expectation Maximization

- If a variable is unobserved or partially observed, you can use all the possible values based on its posterior probability
- The likelihood function can **sum out these variables** and then **update the parameters which maximize the likelihood function**
- Intuitive explanation
 - Given parameters do **INFERENCE** for unobserved data
 - Maximize the marginal likelihood according to inference results (**LEARNING**)
 - Iteratively do inference and learning

General Principles

- The likelihood
 - Maximize the marginal over observed variables
 - $\max_{\theta} P(D_{obs}|\theta) = \max_{\theta} \int P(D_{obs}, X_{miss}|\theta) dX_{miss}$
 - Maximize the MAP over observed variables
 - $\max_{\theta} P(D_{obs}|\theta) \approx \max_{\theta} P(D_{obs}, \hat{X}_{miss}|\theta)$
- Iterative approaches
 - Define a scoring function (for example, likelihood function plus regularization terms)
 - Set parameters \Rightarrow do inference \Rightarrow Re-estimate parameters \Rightarrow re-do inference \Rightarrow ...

GMM Inference Alternative: MAP

- The posterior as above

$$Q(x = k) = P(x = k|y, \theta) = \frac{\pi_k^{(t)} N_k^{(t)}(y[m])}{\sum_{k=1}^K \pi_k^{(t)} N_k^{(t)}(y[m])}$$

$$\tilde{Q}(x[m] = k) = \pi_k^{(t)} N_k^{(t)}(y[m])$$

- Set $x[m]$ as the MAP

$$x[m] = \arg \max_k \tilde{Q}(x[m] = k)$$

- Update the parameters based on the full samples (MLE)

Principle: assign the m -th sample to the most likely component

GMM Inference Alternative: Sampling

- The posterior as above

$$Q(x = k) = P(x = k|y, \theta) = \frac{\pi_k^{(t)} N_k^{(t)}(y[m])}{\sum_{k=1}^K \pi_k^{(t)} N_k^{(t)}(y[m])}$$

$$\tilde{Q}(x[m] = k) = \pi_k^{(t)} N_k^{(t)}(y[m])$$

- Directly sampling from $x \sim \tilde{Q}(x = k)$ to get full samples $(x[m], y[m])$
- Update the parameters based on the full samples (MLE)

Principle: use **samples from the posterior** to approximate the distribution

GMM Learning Alternative: Bayes Learning

- The same for inference

$$\tilde{Q}(x[m] = k) = \pi_k^{(t)} N_k^{(t)}(y[m])$$


- Learning: re-calculate the posterior distributions of parameters based on the complete data

- Priors

- $p(\pi_k) = \text{Dir}(\alpha)$; $p(\mu_k) = N\left(0, \frac{1}{L}\right)$; $p(\Sigma_k) = N(0, I)$

- Posteriors


- $p(\pi_k | \mathcal{D}) = \text{Dirichlet}(\alpha_k + M^*(k))$
 - $p(\mu_k | \mathcal{D}) \propto p(\mathcal{D} | \mu_k) p(\mu_k)$
 - $p(\Sigma_k | \mathcal{D}) \propto p(\mathcal{D} | \Sigma_k) p(\Sigma_k)$



Baum–Welch Algorithm: Revisited

$$\text{Compute } \underset{\theta}{\operatorname{argmax}} P(X|\theta)$$

- General strategy
 - 1) set an initial value of the parameters
 - 2) then do ***inferences of hidden values***
 - 3) then ***re-estimate or re-learn the parameters***
 - 4) Repeat the processes till **convergence**



- Possible problems
 - No guarantee of **convergence**
 - No guarantee of **global optimization**

Baum–Welch Algorithm

Compute $\operatorname{argmax}_{\theta} P(X|\theta)$

- Define an intermediate variables for inference
 - The probability of $y=i$ for time point t and $y=j$ for the following time point $t+1$

$$\xi_t(i, j) = P(y_t = i, y_{t+1} = j | X, \theta)$$

Principle: due to the *Markov* property, all the local transitions are “equal”.

Baum–Welch Algorithm

Compute $\operatorname{argmax}_{\theta} P(X|\theta)$

$$\theta^0 = \{T^0, E^0, \pi^0\}$$

- Use forward and backward algorithm to calculate probability

$$\alpha_t(i) = P(x_1, \dots, x_t = i | \theta^0)$$

$$\beta_t(i) = P(x_{t+1}, \dots, x_T = i | \theta^0)$$

$$\xi_t(i, j) = P(y_t = i, y_{t+1} = j | X, \theta)$$



$$\xi_t(i, j) = \frac{\alpha_t(i) t_{i,j} e_{j, x_{t+1}} \beta_{t+1}(j)}{\sum_{i=1}^Y \sum_{j=1}^Y \alpha_t(i) t_{i,j} \beta_{t+1}(j) e_{j, x_{t+1}}}$$

Inference for the
local transitions

Baum–Welch Algorithm

Compute $\operatorname{argmax}_{\theta} P(X|\theta)$

- The probability of $Y=i$ for time point t :

$$\gamma_t(i) = \sum_{j=1}^Y \xi_t(i, j)$$

- So expected times for state stayed in $Y=i$ and the expected times for state transition $i-j$:

$$\sum_{t=0}^{T-1} \gamma_t(i)$$

$$\sum_{t=0}^{T-1} \xi_t(i, j)$$

Baum–Welch Algorithm

Compute $\operatorname{argmax}_{\theta} P(X|\theta)$

- Re-estimate all parameters (MLE)

$$t_{i,j} = \frac{\sum_{t=0}^{T-1} \xi_t(i,j)}{\sum_{t=0}^{T-1} \gamma_t(i)} \quad e_{i,x} = \frac{\sum_{t=0}^T \mathbf{I}(x_t = x) \gamma_t(i)}{\sum_{t=0}^T \gamma_t(i)}$$

- Repeat above steps until convergence

$$\left| \log(P(X|\theta)) - \log(P(X|\theta^0)) \right| < \varepsilon$$

Generate the hidden variables using current parameter $Y^t \sim P(Y|X, \theta^t)$

Replace the Inference by Sampling

- Initialization: generate the labeling sequencing according to the prior probability

$$\pi \Rightarrow Y_t = y_i, \quad 1 \leq t \leq T$$

- Re-generate Y_t according to the initial setting

$$\begin{aligned} P(y_t | Y, X, \theta) &= P(y_t | y_{t-1}, y_{t+1}, x_t, \theta) \\ &= \frac{P(y_{t+1}, x_t | y_t, y_{t-1}, \theta) P(y_t | y_{t-1}, \theta)}{P(y_{t+1}, x_t | y_{t-1}, \theta)} \end{aligned}$$

We get a “score” proportional to the probability. So we can randomly generate Y_t based on these scores.

$$\propto P(y_{t+1} | y_t, \theta) P(x_t | y_t, \theta) P(y_t | y_{t-1}, \theta)$$

$$= t_{y_t, y_{t+1}} e_{y_t, e_t} t_{y_{t-1}, y_t}$$

MLE for Updating Parameters

- Simply update the parameter in transition and emission probability matrix

$$t_{i \rightarrow j} = \frac{\sum_{t=1}^T I(Y^{t+1}=j, Y^t=i)}{\sum_{t=1}^T I(Y^t=i)}$$


Note: $I(\cdot)$ is an indicator function

$$e_{i \rightarrow k} = \frac{\sum_{t=1}^T I(Y^t=i, X^t=k)}{\sum_{t=1}^T I(Y^t=i)}$$

- *Q: Why does Gibbs Sampling also work?*

Learning in RBMs

3. An RBM (Restricted Boltzmann Machine) is a bipartite Markov network consisting of a visible (observed) layer and a hidden layer, where each node is a 0-1 binary random variable. Consider the following RBM:



Hidden variables

Observed variables

The joint distribution of a configuration is given by:

$$P(H = h, V = v) = \frac{1}{Z} e^{-E(h, v)} \quad h = (h_1, h_2)^T, v = (v_1, \dots, v_5)^T$$

And:

$$E(h, v) = -\sum_{i=1}^2 a_i h_i - \sum_{j=1}^5 b_j v_j - \sum_{ij} w_{ij} h_i v_j$$


Inference step: all parameters are given

$$\tilde{E}(h_1) = -a_1 h_1 - \sum_{j=1}^5 w_{1,j} h_1 v_j$$

$$\tilde{E}(h_2) = -a_2 h_2 - \sum_{j=1}^5 w_{2,j} h_2 v_j$$

Proportional sampling according to above likelihood

3. An RBM (Restricted Boltzmann Machine) is a bipartite Markov network consisting of a visible (observed) layer and a hidden layer, where each node is a 0-1 binary random variable. Consider the following RBM:



Hidden variables

Observed variables

The joint distribution of a configuration is given by:

$$P(H = h, V = v) = \frac{1}{Z} e^{-E(h, v)} \quad h = (h_1, h_2)^T, v = (v_1, \dots, v_5)^T$$

And:

$$E(h, v) = -\sum_{i=1}^2 a_i h_i - \sum_{j=1}^5 b_j v_j - \sum_{ij} w_{ij} h_i v_j$$

Learning step: all variables are observed

Recall: parameter learning in Markov Networks (gradient ascent)


$$\frac{1}{M} l(\theta : x) = a_1 E(h_1) + a_2 E(h_2) + \sum_{j=1}^5 w_{1,j} E(h_1 v_j) + \sum_{j=1}^5 w_{2,j} E(h_2 v_j) + \sum_{j=1}^5 b_j E(v_j) - \ln Z$$

$$\frac{1}{M} \frac{\partial l(\theta : x)}{\partial \theta} \Rightarrow \theta^{t+1} = \theta^t + \lambda^t \frac{1}{M} \frac{\partial l(\theta : x)}{\partial \theta}$$

OR *stochastic gradient ascent*
sample by sample!

How About Other Models?

- Most popular probabilistic graphic models contain hidden variables:
 - For latent factor analysis?
 - For restricted Boltzmann machine?
 - For latent *Dirichlet* allocation?
 - For conditional random fields?
 - ...



$$x: \alpha_1 = \alpha_2 = \alpha_3 = \alpha$$


$$y: \beta = 2\alpha$$

$$xx: w_{12} = w_{13} = w_{23} = w$$

$$xy: h_{1y} = h_{2y} = h_{3y} = -w$$

Q: If Y is a hidden variable, how to design a learning algorithm?

Inference Step



$$x: \alpha_1 = \alpha_2 = \alpha_3 = \alpha$$

$$y: \beta = 2\alpha$$

$$xx: w_{12} = w_{13} = w_{23} = w$$

$$xy: h_{1y} = h_{2y} = h_{3y} = -w$$

$$P(x, y | \alpha, w) = \frac{1}{Z} \exp \left\{ \alpha (x_1 + x_2 + x_3 + 2y) + w (x_1 x_2 + x_2 x_3 + x_1 x_3 - y (x_1 + x_2 + x_3)) \right\}$$

$$P(y | x, \alpha, w) \propto P(y, x | \alpha, w)$$

$$\Rightarrow \frac{P(y=1, x | \alpha, w)}{P(y=0, x | \alpha, w)} = \exp \left\{ 2\alpha + w (-(x_1 + x_2 + x_3)) \right\} = s$$

$$\Rightarrow P(y=1, x | \alpha, w) = \frac{s}{1+s}, \quad P(y=0, x | \alpha, w) = \frac{1}{1+s}$$

data-hw7.csv			
	x1	x2	x3
1	0	1	1
2	0	0	1
3	1	0	1
4	0	1	0
5	1	1	0
6	0	1	1
7	1	1	1
8	1	0	0
9	0	1	1
10	1	0	1
11	1	0	1

You can either use **the expectation (probability)** or **sampling** to fill the single hidden variable!!

Learning Step (SGD)

$$P(x, y | \alpha, w) = \frac{1}{Z} \exp \left\{ \alpha (x_1 + x_2 + x_3 + 2y) + w (x_1 x_2 + x_2 x_3 + x_1 x_3 - y (x_1 + x_2 + x_3)) \right\}$$

- SGD: $\boldsymbol{\theta}^{(k+1)} = \boldsymbol{\theta}^{(k)} + \lambda^{(k)} \nabla l_{m_k}(\boldsymbol{\theta}; \mathbf{x}[m_k])$
- Calculate the partial derivatives
 - For α
 - For w

$$\begin{aligned} \frac{\partial}{\partial \theta_i} \ln Z(\boldsymbol{\theta}) &= \frac{1}{Z(\boldsymbol{\theta})} \sum_{\xi} \frac{\partial}{\partial \theta_i} \exp \left\{ \sum_j \theta_j f_j(\xi) \right\} \\ &= \frac{1}{Z(\boldsymbol{\theta})} \sum_{\xi} f_i(\xi) \exp \left\{ \sum_j \theta_j f_j(\xi) \right\} \\ &= E_{\boldsymbol{\theta}}[f_i]. \end{aligned}$$

How to Deal With Missing Values?


- Under the framework of probabilistic inferences, just use the same strategy as filling the hidden variables!

Generalization



- Inference
 - MCMC
 - Variational inference
 - Belief propagation
- Learning
 - MLE, MAP
 - Gradient ascent/descent
 - Hierarchical Bayesian

The Markov blanket in biotic systems



Free energy and active inference

Karl Friston, University College London

lemma: any (ergodic random) dynamical system (m) that possesses a Markov blanket will appear to actively maintain its structural and dynamical integrity

\vdots

ω



$p(x|m)$

The Fokker-Planck equation \vdots

$$\nabla \cdot (\Gamma \nabla - f) p$$


And its solution in terms of curl-free and divergence-free components

\vdots

$$0 \Leftrightarrow f(x) = (\Gamma - Q) \nabla \ln p(x|m)$$



But what about the Markov blanket?




$$f_{\mu}(s, a, \mu) = (\Gamma - Q) \nabla_{\mu} \ln p(s, a, \mu | m)$$

$$f_a(s, a, \mu) = (\Gamma - Q) \nabla_a \ln p(s, a, \mu | m)$$

$$\ln p(s, a, \mu | m) = \text{Value}$$



Reinforcement learning, optimal control
and expected utility theory



Pavlov

$$-\ln p(s, a, \mu | m) = \text{Surprise (free energy)}$$



Information theory and minimum
redundancy



Barlow

$$E_t[-\ln p(s, a, \mu | m)] = \text{Entropy}$$



Self-organisation, cybernetics and
homoeostasis



Ashby

$$p(\tilde{\gamma} \sim \tilde{\gamma})$$

$$= \text{Model evidence}$$

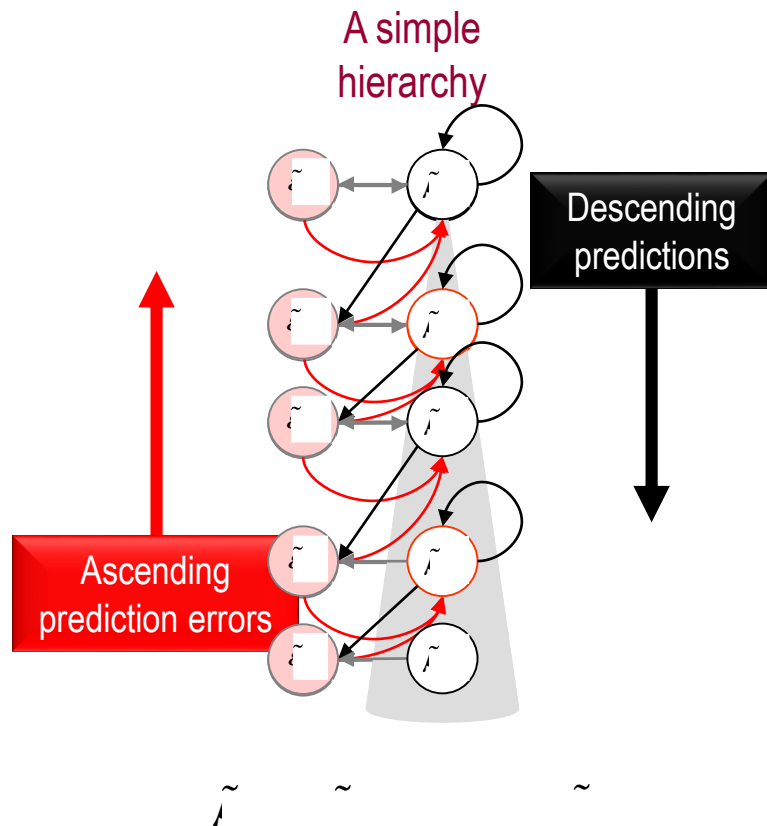


Bayesian brain, active inference and
predictive coding




Helmholtz

Hierarchical generative models

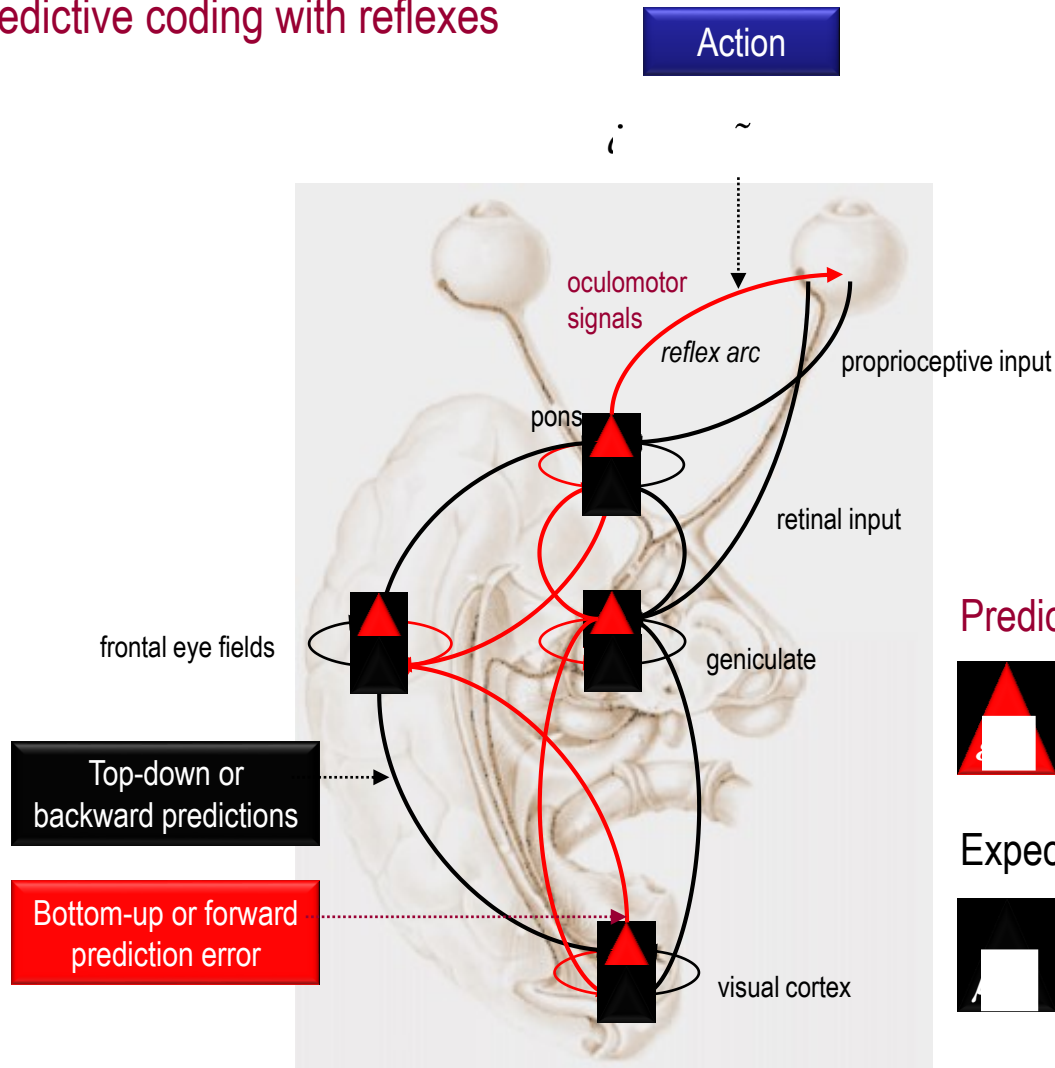


Kalman filter





Predictive coding with reflexes



Perception

Prediction error (superficial pyramidal cells)



Expectations (deep pyramidal cells)



E-M algorithm!

Message passing / Belief propagation

Why Need Deep Latent Models?

- Low-order models cannot fully understand the data generated from high-order models

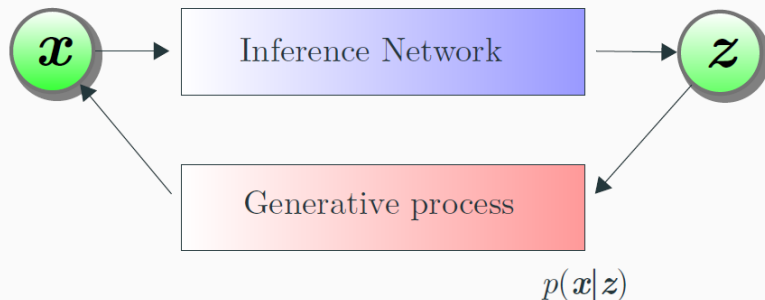


Deep VAE & GAN

- VAE
 - Target distribution in *low-dimension subspace*
 - Minimize the likelihood

$$\begin{aligned}\mathcal{L}(\theta; \mathbf{x}) &= \mathbb{E}_{z \sim q(z|\mathbf{x}; \theta)} \log p(\mathbf{x}, z) + \mathbb{H}(q(z|\mathbf{x}; \theta)) \\ &= \mathbb{E}_{z \sim q(z|\mathbf{x}; \theta)} \log p(\mathbf{x}|z) - \mathbb{KL}(q(z|\mathbf{x}; \theta) \parallel p(z))\end{aligned}$$


$$q(z|\mathbf{x}) = \mathcal{N}(f_{\mu}(\mathbf{x}), f_{\sigma}(\mathbf{x}))$$



- GAN
 - Learn a *latent generator* by mapping noise to data
 - Minimize the discrimination

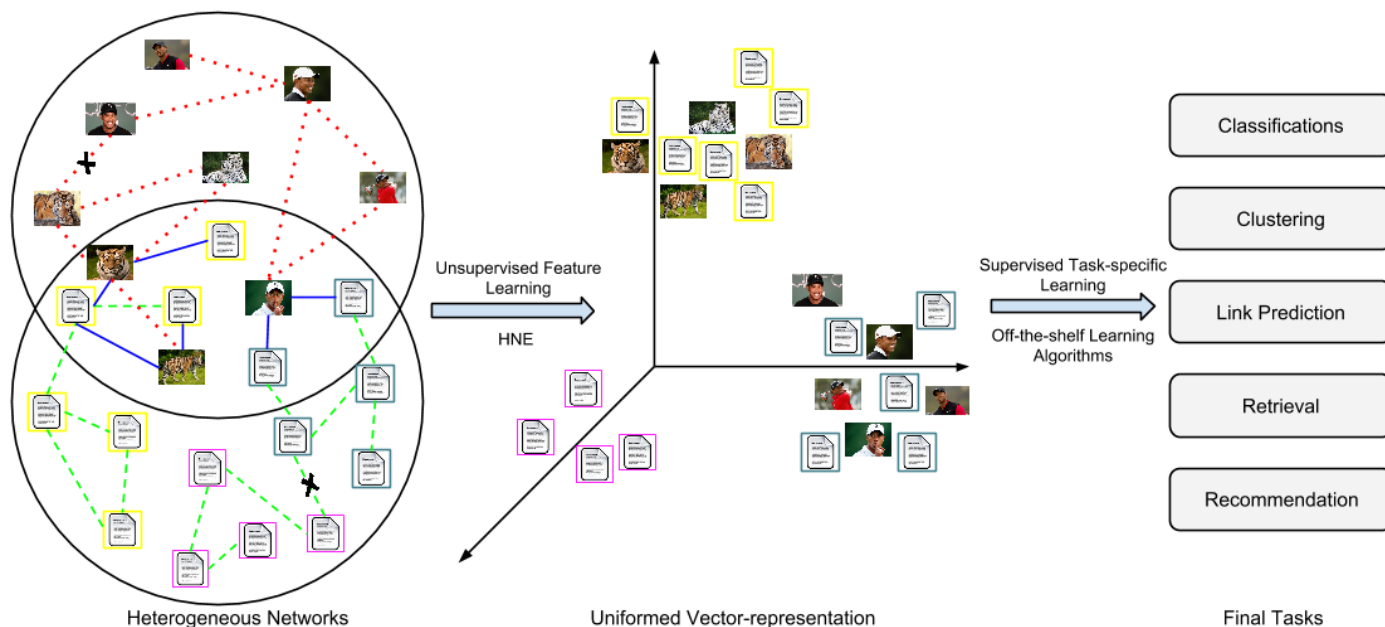
$$\min_{\mathcal{G}} \max_{\mathcal{D}} V(\mathcal{D}, \mathcal{G}) =$$

$$\mathbb{E}_{\mathbf{x} \sim p_{data}(\mathbf{x})} [\log \mathcal{D}(\mathbf{x})] + \mathbb{E}_{(\mathbf{z}) \sim p_z((\mathbf{z}))} [\log (1 - \mathcal{D}(\mathcal{G}(\mathbf{z})))]$$




Embedding for Relevant Networks

- Task: mapping nodes in networks to a latent Euclidean space (low-dimensional subspace) by preserving their “similarities”



LINE: Second-Order Similarity



Second-order relations are more **stable** than the *first-order*

- First-order similarity

$$p_1(v_i, v_j) = \frac{1}{1 + \exp(-\vec{u}_i^T \cdot \vec{u}_j)}$$

$$O_1 = - \sum_{(i,j) \in E} w_{ij} \log p_1(v_i, v_j)$$

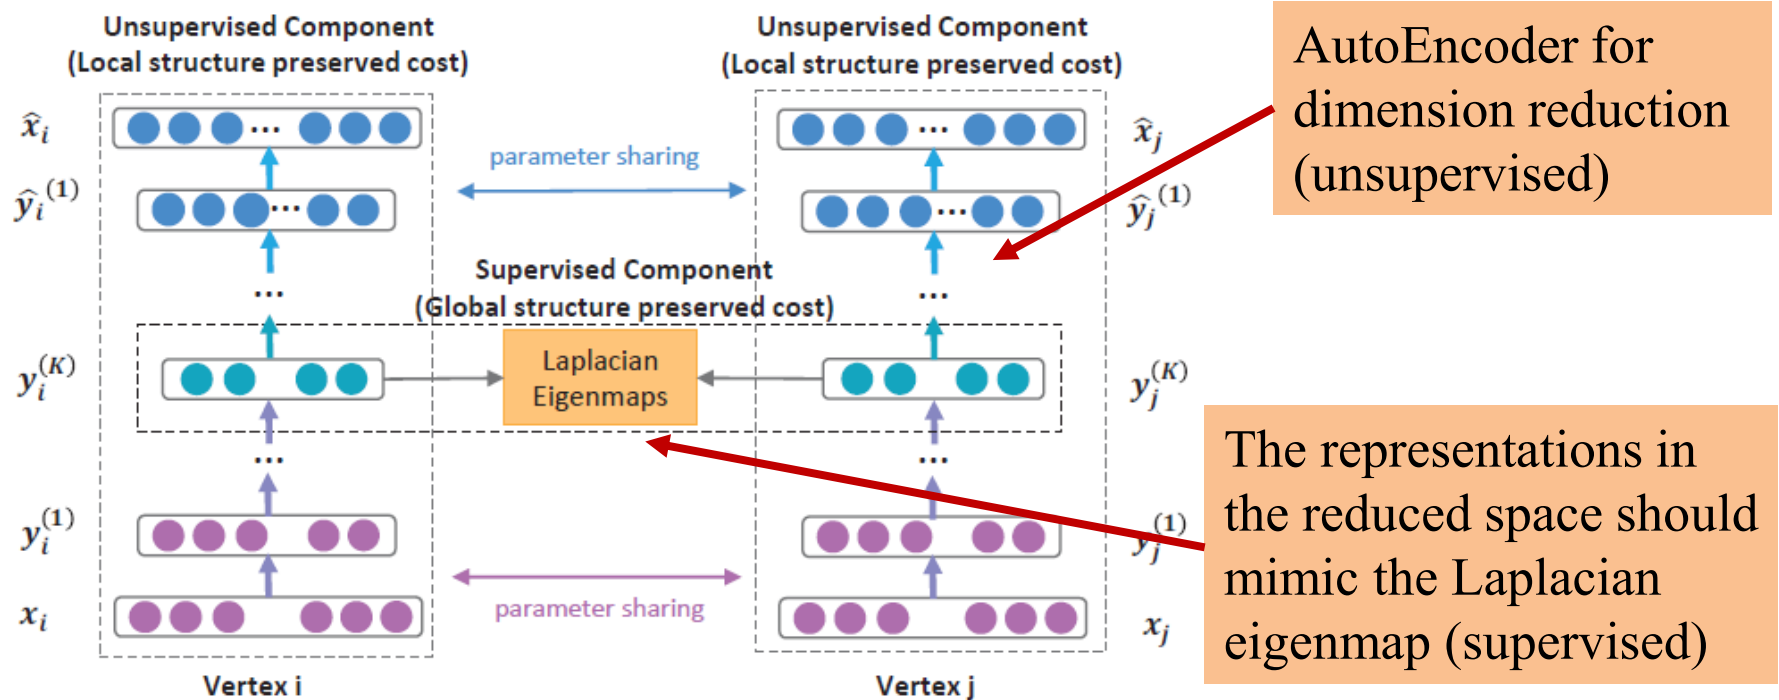
- Second-order similarity

$$p_2(v_j | v_i) = \frac{\exp(\vec{u}_j'^T \cdot \vec{u}_i)}{\sum_{k=1}^{|V|} \exp(\vec{u}_k'^T \cdot \vec{u}_i)}$$

$$O_2 = - \sum_{(i,j) \in E} w_{ij} \log p_2(v_j | v_i)$$

Deep Structural Network Embedding

- Aim: find a low-dimensional representation of the adjacent matrix S



Deep Structural Network Embedding

- First-order penalty

$$\begin{aligned}\mathcal{L}_{1st} &= \sum_{i,j=1}^n s_{i,j} \|\mathbf{y}_i^{(K)} - \mathbf{y}_j^{(K)}\|_2^2 \\ &= \sum_{i,j=1}^n s_{i,j} \|\mathbf{y}_i - \mathbf{y}_j\|_2^2\end{aligned}$$

- Second-order penalty

$$\begin{aligned}\mathcal{L}_{2nd} &= \sum_{i=1}^n \|(\hat{\mathbf{x}}_i - \mathbf{x}_i) \odot \mathbf{b}_i\|_2^2 \\ &= \|(\hat{X} - X) \odot B\|_F^2\end{aligned}$$

$$\begin{aligned}\mathcal{L}_{mix} &= \mathcal{L}_{2nd} + \alpha \mathcal{L}_{1st} + \nu \mathcal{L}_{reg} \\ &= \|(\hat{X} - X) \odot B\|_F^2 + \alpha \sum_{i,j=1}^n s_{i,j} \|\mathbf{y}_i - \mathbf{y}_j\|_2^2 + \nu \mathcal{L}_{reg}\end{aligned}$$

$$\mathcal{L}_{reg} = \frac{1}{2} \sum_{k=1}^K (\|W^{(k)}\|_F^2 + \|\hat{W}^{(k)}\|_F^2)$$



$$\begin{aligned}\frac{\partial \mathcal{L}_{mix}}{\partial \hat{W}^{(k)}} &= \frac{\partial \mathcal{L}_{2nd}}{\partial \hat{W}^{(k)}} + \nu \frac{\partial \mathcal{L}_{reg}}{\partial \hat{W}^{(k)}} \\ \frac{\partial \mathcal{L}_{mix}}{\partial W^{(k)}} &= \frac{\partial \mathcal{L}_{2nd}}{\partial W^{(k)}} + \alpha \frac{\partial \mathcal{L}_{1st}}{\partial W^{(k)}} + \nu \frac{\partial \mathcal{L}_{reg}}{\partial W^{(k)}}, k = 1, \dots, K\end{aligned}$$

Deep Structural Network Embedding

Algorithm 1 Training Algorithm for the semi-supervised deep model of *SDNE*

Input: the network $G = (V, E)$ with adjacency matrix S , the parameters α and ν

Output: Network representations Y and updated Parameters: θ

1: Pretrain the model through deep belief network to obtain the initialized parameters $\theta = \{\theta^{(1)}, \dots, \theta^{(K)}\}$

2: $X = S$

3: **repeat**

4: Based on X and θ , apply Eq. 1 to obtain \hat{X} and $Y = Y^K$. *Inference*

5: $\mathcal{L}_{mix}(X; \theta) = \|(\hat{X} - X) \odot B\|_F^2 + 2\alpha \text{tr}(Y^T L Y) + \nu \mathcal{L}_{reg}.$

6: Based on Eq. 6, use $\partial \mathcal{L}_{mix} / \partial \theta$ to back-propagate through the entire network to get updated parameters θ . *Learning*

7: **until** converge

8: Obtain the network representations $Y = Y^{(K)}$

Summary: General Principles

- The likelihood
 - Maximize the marginal over observed variables
 - $\max_{\theta} P(D_{obs}|\theta) = \max_{\theta} \int P(D_{obs}, X_{miss}|\theta) dX_{miss}$
 - Maximize the MAP over observed variables
 - $\max_{\theta} P(D_{obs}|\theta) \approx \max_{\theta} P(D_{obs}, \hat{X}_{miss}|\theta)$
- Iterative approaches
 - Define a scoring function (for example, likelihood function plus regularization terms)
 - Set parameters => do inference => Re-estimate parameters
=> re-do inference => ...

General Principles Extended

- Define an objective/scoring function
 - With both parameters and hidden variables (or missing values)
 - Set parameters \Rightarrow do inference \Rightarrow re-estimate parameters \Rightarrow re-do inference \Rightarrow ...
 - Stop until the objective/scoring function converges (or the posterior probability converges for Bayesian models)

Chapter 11 The End

Do Inference & Learning

Alternatively